

Assignment3

Task 1

[acadgild@localhost mapreduce]\$ hdfs dfs -cat hadoop/word-count.txt

Java HotSpot(TM) Client VM warning: You have loaded library </home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0> which might have disabled stack guard. The VM will try to fix the stack guard now. It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/06/09 17:59:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

GNU nano 2.0.9 File: word-count.txt Modified

i love hadoop

[New File]

^G Get Help ^O WriteOut ^R Read File ^Y Prev Page ^K Cut Text ^C Cur Pos

^X Exit^J Justify ^W Where Is ^V Next Page ^U UnCut Text^T To Spell

Wordcount

[acadgild@localhost mapreduce]\$ hadoop jar hadoop-mapreduce-examples-2.7.2.jar wordcount hadoop/word-count.txt hadoop/word-countout1

Java HotSpot(TM) Client VM warning: You have loaded library </home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0> which might have disabled stack guard. The VM will try to fix the stack guard now. It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/06/09 18:05:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/06/09 18:05:16 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/06/09 18:05:17 INFO input.FileInputFormat: Total input paths to process : 1

18/06/09 18:05:17 INFO mapreduce.JobSubmitter: number of splits:1

18/06/09 18:05:18 INFO mapreduce.JobSubmitter: Submitting tokens for job:

job_1528539845120_0004

18/06/09 18:05:18 INFO impl.YarnClientImpl: Submitted application application_1528539845120_0004

18/06/09 18:05:19 INFO mapreduce.Job: The url to track the job: [http://localhost:8088/proxy/](http://localhost:8088/proxy/application_1528539845120_0004/)

[application_1528539845120_0004/](http://localhost:8088/proxy/application_1528539845120_0004/)

18/06/09 18:05:19 INFO mapreduce.Job: Running job: job_1528539845120_0004

18/06/09 18:05:33 INFO mapreduce.Job: Job job_1528539845120_0004 running in uber mode : false

18/06/09 18:05:33 INFO mapreduce.Job: map 0% reduce 0%

18/06/09 18:05:42 INFO mapreduce.Job: map 100% reduce 0%

18/06/09 18:05:52 INFO mapreduce.Job: map 100% reduce 100%

18/06/09 18:05:52 INFO mapreduce.Job: Job job_1528539845120_0004 completed successfully

18/06/09 18:05:53 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=698

FILE: Number of bytes written=236741

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=2066

HDFS: Number of bytes written=521

HDFS: Number of read operations=6

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=6776
Total time spent by all reduces in occupied slots (ms)=7144
Total time spent by all map tasks (ms)=6776
Total time spent by all reduce tasks (ms)=7144
Total vcore-milliseconds taken by all map tasks=6776
Total vcore-milliseconds taken by all reduce tasks=7144
Total megabyte-milliseconds taken by all map tasks=6938624
Total megabyte-milliseconds taken by all reduce tasks=7315456

Map-Reduce Framework

Map input records=24
Map output records=63
Map output bytes=2154
Map output materialized bytes=698
Input split bytes=122
Combine input records=63
Combine output records=43
Reduce input groups=43
Reduce shuffle bytes=698
Reduce input records=43
Reduce output records=43
Spilled Records=86
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=234
CPU time spent (ms)=1790
Physical memory (bytes) snapshot=229429248
Virtual memory (bytes) snapshot=636829696
Total committed heap usage (bytes)=137498624

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1944

File Output Format Counters

Bytes Written=521

Wordmedian

```
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.7.2.jar  
wordmedian hadoop/word-count.txt hadoop/word-medianout
```

Java HotSpot(TM) Client VM warning: You have loaded library </home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0> which might have disabled stack guard. The VM will try to fix the stack guard now. It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/06/09 18:00:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/06/09 18:00:23 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/06/09 18:00:24 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not

performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

18/06/09 18:00:25 INFO input.FileInputFormat: Total input paths to process : 1

18/06/09 18:00:25 INFO mapreduce.JobSubmitter: number of splits:1

18/06/09 18:00:26 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1528539845120_0003

18/06/09 18:00:26 INFO impl.YarnClientImpl: Submitted application application_1528539845120_0003

18/06/09 18:00:26 INFO mapreduce.Job: The url to track the job: [http://localhost:8088/proxy/
application_1528539845120_0003/](http://localhost:8088/proxy/application_1528539845120_0003/)

18/06/09 18:00:26 INFO mapreduce.Job: Running job: job_1528539845120_0003

18/06/09 18:00:42 INFO mapreduce.Job: Job job_1528539845120_0003 running in uber mode : false

18/06/09 18:00:42 INFO mapreduce.Job: map 0% reduce 0%

18/06/09 18:00:51 INFO mapreduce.Job: map 100% reduce 0%

18/06/09 18:01:02 INFO mapreduce.Job: map 100% reduce 100%

18/06/09 18:01:03 INFO mapreduce.Job: Job job_1528539845120_0003 completed successfully

18/06/09 18:01:04 INFO mapreduce.Job: Counters: 49

File System Counters

- FILE: Number of bytes read=146
- FILE: Number of bytes written=235369
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=2066
- HDFS: Number of bytes written=64
- HDFS: Number of read operations=6
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=2

Job Counters

- Launched map tasks=1
- Launched reduce tasks=1
- Data-local map tasks=1
- Total time spent by all maps in occupied slots (ms)=7700
- Total time spent by all reduces in occupied slots (ms)=7208
- Total time spent by all map tasks (ms)=7700
- Total time spent by all reduce tasks (ms)=7208
- Total vcore-milliseconds taken by all map tasks=7700
- Total vcore-milliseconds taken by all reduce tasks=7208
- Total megabyte-milliseconds taken by all map tasks=7884800
- Total megabyte-milliseconds taken by all reduce tasks=7380992

Map-Reduce Framework

- Map input records=24
- Map output records=63
- Map output bytes=504
- Map output materialized bytes=146
- Input split bytes=122
- Combine input records=63
- Combine output records=14
- Reduce input groups=14
- Reduce shuffle bytes=146
- Reduce input records=14
- Reduce output records=14
- Spilled Records=28
- Shuffled Maps =1
- Failed Shuffles=0
- Merged Map outputs=1
- GC time elapsed (ms)=240
- CPU time spent (ms)=2090
- Physical memory (bytes) snapshot=229339136

Virtual memory (bytes) snapshot=635781120
Total committed heap usage (bytes)=137498624
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=1944
File Output Format Counters
Bytes Written=64

The median is: 5

wordmean

[acadgild@localhost mapreduce]\$ hadoop jar hadoop-mapreduce-examples-2.7.2.jar wordmean hadoop/word-count.txt hadoop/word-mean

Java HotSpot(TM) Client VM warning: You have loaded library </home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0> which might have disabled stack guard. The VM will try to fix the stack guard now. It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/06/09 18:06:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/06/09 18:06:16 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/06/09 18:06:18 INFO input.FileInputFormat: Total input paths to process : 1

18/06/09 18:06:18 INFO mapreduce.JobSubmitter: number of splits:1

18/06/09 18:06:18 INFO mapreduce.JobSubmitter: Submitting tokens for job:

job_1528539845120_0005

18/06/09 18:06:19 INFO impl.YarnClientImpl: Submitted application application_1528539845120_0005

18/06/09 18:06:19 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1528539845120_0005/

18/06/09 18:06:19 INFO mapreduce.Job: Running job: job_1528539845120_0005

18/06/09 18:06:33 INFO mapreduce.Job: Job job_1528539845120_0005 running in uber mode : false

18/06/09 18:06:33 INFO mapreduce.Job: map 0% reduce 0%

18/06/09 18:06:42 INFO mapreduce.Job: map 100% reduce 0%

18/06/09 18:06:53 INFO mapreduce.Job: map 100% reduce 100%

18/06/09 18:06:54 INFO mapreduce.Job: Job job_1528539845120_0005 completed successfully

18/06/09 18:06:54 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=39

FILE: Number of bytes written=235413

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=2066

HDFS: Number of bytes written=21

HDFS: Number of read operations=6

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

Job Counters

Launched map tasks=1

Launched reduce tasks=1

Data-local map tasks=1

Total time spent by all maps in occupied slots (ms)=6653

Total time spent by all reduces in occupied slots (ms)=6796
 Total time spent by all map tasks (ms)=6653
 Total time spent by all reduce tasks (ms)=6796
 Total vcore-milliseconds taken by all map tasks=6653
 Total vcore-milliseconds taken by all reduce tasks=6796
 Total megabyte-milliseconds taken by all map tasks=6812672
 Total megabyte-milliseconds taken by all reduce tasks=6959104
 Map-Reduce Framework
 Map input records=24
 Map output records=126
 Map output bytes=1827
 Map output materialized bytes=39
 Input split bytes=122
 Combine input records=126
 Combine output records=2
 Reduce input groups=2
 Reduce shuffle bytes=39
 Reduce input records=2
 Reduce output records=2
 Spilled Records=4
 Shuffled Maps =1
 Failed Shuffles=0
 Merged Map outputs=1
 GC time elapsed (ms)=257
 CPU time spent (ms)=1770
 Physical memory (bytes) snapshot=229376000
 Virtual memory (bytes) snapshot=636829696
 Total committed heap usage (bytes)=137498624
 Shuffle Errors
 BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0
 File Input Format Counters
 Bytes Read=1944
 File Output Format Counters
 Bytes Written=21

The mean is: 29.19047619047619

word standarddeviation

[acadgild@localhost mapreduce]\$ hadoop jar hadoop-mapreduce-examples-2.7.2.jar wordstandarddeviation hadoop/word-count.txt hadoop/word-standarddeviation

Java HotSpot(TM) Client VM warning: You have loaded library </home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0> which might have disabled stack guard. The VM will try to fix the stack guard now. It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/06/09 18:09:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/06/09 18:09:13 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/06/09 18:09:15 INFO input.FileInputFormat: Total input paths to process : 1

18/06/09 18:09:15 INFO mapreduce.JobSubmitter: number of splits:1

18/06/09 18:09:15 INFO mapreduce.JobSubmitter: Submitting tokens for job:

job_1528539845120_0006

18/06/09 18:09:16 INFO impl.YarnClientImpl: Submitted application application_1528539845120_0006

18/06/09 18:09:16 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1528539845120_0006/
18/06/09 18:09:16 INFO mapreduce.Job: Running job: job_1528539845120_0006
18/06/09 18:09:30 INFO mapreduce.Job: Job job_1528539845120_0006 running in uber mode : false
18/06/09 18:09:30 INFO mapreduce.Job: map 0% reduce 0%
18/06/09 18:09:40 INFO mapreduce.Job: map 100% reduce 0%
18/06/09 18:09:50 INFO mapreduce.Job: map 100% reduce 100%
18/06/09 18:09:50 INFO mapreduce.Job: Job job_1528539845120_0006 completed successfully
18/06/09 18:09:50 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=56
FILE: Number of bytes written=235637
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=2066
HDFS: Number of bytes written=35
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=6948
Total time spent by all reduces in occupied slots (ms)=7015
Total time spent by all map tasks (ms)=6948
Total time spent by all reduce tasks (ms)=7015
Total vcore-milliseconds taken by all map tasks=6948
Total vcore-milliseconds taken by all reduce tasks=7015
Total megabyte-milliseconds taken by all map tasks=7114752
Total megabyte-milliseconds taken by all reduce tasks=7183360

Map-Reduce Framework

Map input records=24
Map output records=189
Map output bytes=2772
Map output materialized bytes=56
Input split bytes=122
Combine input records=189
Combine output records=3
Reduce input groups=3
Reduce shuffle bytes=56
Reduce input records=3
Reduce output records=3
Spilled Records=6
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=241
CPU time spent (ms)=1900
Physical memory (bytes) snapshot=229335040
Virtual memory (bytes) snapshot=635932672
Total committed heap usage (bytes)=137498624

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0

WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=1944
File Output Format Counters
Bytes Written=35

The standard deviation is: 34.907068071674885